

Topics in Probabilistic Modeling and Inference  
(CS698X, Spring 2019)  
Homework 1  
Due Date: Feb 8, 2019 (11:59pm)

**Instructions:**

- Only electronic submissions will be accepted. Your main PDF writeup must be typeset in LaTeX (please also refer to the “Additional Instructions” below).
- Your submission will have two parts: The main PDF writeup (to be submitted via Gradescope <https://www.gradescope.com/>) and the code for the programming part (to be submitted via this Dropbox link: <https://tinyurl.com/cs698x-sp18-hw1-code>). Both parts must be submitted by the deadline to receive full credit (**delay in submitting either part would incur late penalty for both parts**). We will be accepting late submissions upto 72 hours after the deadline (with every 24 hours delay incurring a 10% late penalty). We won’t be able to accept submissions after that.
- We have created your Gradescope account (you should have received the notification). Please use your IITK CC ID (not any other email ID) to login. Use the “Forgot Password” option to set your password.

**Additional Instructions**

- We have provided a LaTeX template file `hw1sol.tex` to help typeset your PDF writeup. There is also a style file `pmi.sty` that contain shortcuts to many of the useful LaTeX commands for doing things such as boldfaced/calligraphic fonts for letters, various mathematical/greek symbols, etc., and others. Use of these shortcuts is recommended (but not necessary).
- Your answer to every question should begin on a new page. The provided template is designed to do this automatically. However, if it fails to do so, use the `\clearpage` option in LaTeX before starting the answer to a new question, to *enforce* this.
- While submitting your assignment on the Gradescope website, you will have to specify on which page(s) is question 1 answered, on which page(s) is question 2 answered etc. To do this properly, first ensure that the answer to each question starts on a different page.
- Be careful to flush all your floats (figures, tables) corresponding to question  $n$  before starting the answer to question  $n + 1$  otherwise, while grading, we might miss your important parts of your answers.
- Your solutions must appear in proper order in the PDF file i.e. solution to question  $n$  must be complete in the PDF file (including all plots, tables, proofs etc) before you present a solution to question  $n + 1$ .
- For the programming part, all the code and README should be zipped together and submitted as a single file named `yourrollnumber.zip`. Please DO NOT submit the data provided.

## Problem 1 (5 marks)

**(MLE as KL Minimization)** Suppose you are given  $N$  observations  $\{x_1, x_2, \dots, x_N\}$  from some true underlying data distribution  $p_{data}(x)$ . To learn it, you assume a parametrized distribution  $p(x|\theta)$  and estimate the parameters  $\theta$  using MLE. Show that doing MLE is equivalent to finding  $\theta$  that minimizes the KL divergence between the true distribution  $p_{data}(x)$  and the assumed distribution  $p(x|\theta)$ . Note that KL divergence between two probability distributions  $p$  and  $q$  is asymmetric and can be defined in two different ways:  $KL(p||q)$  or  $KL(q||p)$ . For this problem, minimizing only one of these two will be equivalent to MLE. Why not the other one?

## Problem 2 (5 marks)

**(Distribution of Empirical Mean of Gaussian Observations)** Consider  $N$  scalar-valued observations  $x_1, \dots, x_N$  drawn i.i.d. from  $\mathcal{N}(\mu, \sigma^2)$ . Consider their empirical mean  $\bar{x} = \frac{1}{N} \sum_{n=1}^N x_n$ . Representing the empirical mean as a linear transformation of a random variable, derive the probability distribution of  $\bar{x}$ .

## Problem 3 (15 marks)

**(Benefits of Hierarchical Modeling?)** Consider a dataset of test-scores of students from  $M$  schools in a district:  $\mathbf{x} = \{\mathbf{x}^{(m)}\}_{m=1}^M = \{x_1^{(m)}, \dots, x_{N_m}^{(m)}\}_{m=1}^M$ , where  $N_m$  denotes the number of students in school  $m$ . Assume the scores of students in school  $m$  to be Gaussian distributed  $x_n^{(m)} \sim \mathcal{N}(\mu_m, \sigma^2)$  where the Gaussian's mean  $\mu_m$  is unknown and the variance  $\sigma^2$  is same for all schools and known (for simplicity). Assume the means  $\mu_1, \dots, \mu_M$  of the  $M$  Gaussians to also be Gaussian distributed  $\mu_m \sim \mathcal{N}(\mu_0, \sigma_0^2)$  where  $\mu_0$  and  $\sigma_0^2$  are hyperparameters.

1. Assume the hyperparameters  $\mu_0$  and  $\sigma_0^2$  to be known. Derive the posterior distribution of  $\mu_m$  and write down the mean and variance of this posterior distribution. **Note:** While you can derive it the usual way, the derivation will be much more compact if you use the result of Problem 2 and think of each school's data as a *single* observation (the empirical mean of observations) having the distribution derived in Problem 2.
2. Assume the hyperparameter  $\mu_0$  to be unknown (but still keep  $\sigma_0^2$  as fixed for simplicity). Derive the marginal likelihood  $p(\mathbf{x}|\mu_0, \sigma^2, \sigma_0^2)$  and use MLE-II to estimate  $\mu_0$  (note again that  $\sigma^2$  and  $\sigma_0^2$  are known here). Note: Looking at the form/expression of the marginal likelihood, if the MLE-II result looks obvious to you, you may skip the derivation and directly write the result.
3. Consider using this MLE-II estimate of  $\mu_0$  from part (2) in the posteriors of each  $\mu_m$  you derived in part (1). Do you see any benefit in using the MLE-II estimate of  $\mu_0$  as opposed to using a known value of  $\mu_0$ ?

## Problem 4 (20 marks)

**Binary Latent Matrices** Consider modeling an  $N \times K$  binary matrix  $\mathbf{Z}$  with its entries assumed to be generated independently as follows

$$\begin{aligned} Z_{nk} | \pi_k &\sim \text{Bernoulli}(\pi_k) & n = 1, \dots, N, k = 1, \dots, K \\ \pi_k &\sim \text{Beta}(\alpha/K, 1) & k = 1, \dots, K \end{aligned}$$

- Integrate out  $\{\pi_k\}_{k=1}^K$  and derive the expression for the marginal prior  $p(\mathbf{Z}|\alpha)$  and show that it can be written in form of a product of ratios of Beta functions.
- Derive the distribution  $p(Z_{nk} | Z_{-nk})$  where  $Z_{-nk}$  denotes all the entries in  $k$ -th column of  $\mathbf{Z}$ , except  $Z_{nk}$ . Since  $Z_{nk}$  is binary, it suffices to compute  $p(Z_{nk} = 1 | Z_{-nk})$  (hint: Use Bayes rule). Explain why the form of the result makes intuitive sense.
- As a function of  $\alpha$ , what will be the expected number of ones in each column of  $\mathbf{Z}$ , and in all of  $\mathbf{Z}$ ?

## Problem 5 (30 marks)

**(Spike-and-Slab Model for Sparsity)** Suppose  $w$  is a real-valued r.v. that can either be close to zero with probability  $\pi$ , or take a wide range of real values with probability  $(1 - \pi)$ . An example of this could be in a regression problem where  $w$  is the weight of some feature. The feature could be irrelevant for predicting the output (in which case we would expect  $w$  to be close to zero) or be useful (in which case we would expect  $w$  to be non-zero with a wide range of possible values). We want to infer  $w$  from data taking a Bayesian approach. Note that, in practice,  $w$  is a vector (with each entry modeled this way) but here we will consider the scalar  $w$  case.

A popular approach to solve such problems is to impose a *spike and slab prior* on  $w$ . Let  $b \in \{0, 1\}$  be a binary random variable and define the following *conditional* prior on  $w$ :

$$p(w|b, \sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2) = \begin{cases} \mathcal{N}(w|0, \sigma_{\text{spike}}^2) & b = 0 \\ \mathcal{N}(w|0, \sigma_{\text{slab}}^2) & b = 1, \end{cases}$$

Depending on the value of  $b$  (which itself is unknown),  $w$  is assumed drawn from one of the two distributions: a “peaky” one  $\mathcal{N}(w|0, \sigma_{\text{spike}}^2)$  with variance  $\sigma_{\text{spike}}^2$  being very small, and a “flat” one  $\mathcal{N}(w|0, \sigma_{\text{slab}}^2)$ , with  $\sigma_{\text{slab}} \gg \sigma_{\text{spike}}$ . So, basically, the value of the binary “mask”  $b$  decides whether the feature is relevant or not.

We usually don’t know  $b$ , so we must either infer it with  $w$ , or marginalize it if we care about the value of  $w$ .

- Assume a prior  $p(b = 1) = \pi = 1/2$ , which means both Gaussians are equally likely for  $w$ . What is the *marginal* prior  $p(w|\sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2)$ , i.e., the prior over  $w$  after integrating out  $b$ ?
- Plot this marginal prior distribution for  $(\sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2) = (1, 100)$ . Briefly comment on how the shape of this distribution compares with that of a typical Gaussian distribution?
- Suppose someone gave us a “noisy” version of  $w$  defined as  $x = w + \epsilon$  where  $\epsilon \sim \mathcal{N}(\epsilon|0, \rho^2)$ . This is equivalent to writing  $p(x|w, \rho^2) = \mathcal{N}(x|w, \rho^2)$ . Assume the variance  $\rho^2$  to be known. Given  $x$ , what is the posterior distribution of  $b$ ,  $p(b = 1|x, \sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2, \rho^2)$ ? Note that  $w$  must NOT appear in this expression (has to be integrated out first). Plot the resulting posterior  $p(b = 1|x, \sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2, \rho^2)$  as a function of  $x$ .
- Given the noisy observation  $x = w + \epsilon$  as defined above, what is the posterior distribution of  $w$ , i.e.,  $p(w|x, \sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2, \rho^2)$ ? Note that  $b$  must NOT appear in this expression (has to be integrated out or summed over since  $b$  is discrete).
- Assume  $(\sigma_{\text{spike}}^2, \sigma_{\text{slab}}^2) = (1, 100)$ , the noise variance  $\rho^2 = 0.01$ . For these settings of the hyperparameters, plot the posterior distribution of  $w$  given a noisy observation  $x = 3$ .

Do not submit the code for this part. All of the answers/derivations for this part (including the plots) should be in the PDF writeup.

## Problem 6 (25 marks): Programming Assignment

**(Bayesian Linear Regression)** Consider a toy data set consisting of 10 training examples  $\{x_n, y_n\}_{n=1}^{10}$  with each input  $x_n$  as well as the output  $y_n$  being scalars. The data is given below.

$$\begin{aligned} \mathbf{x} &= [-2.23, -1.30, -0.42, 0.30, 0.33, 0.52, 0.87, 1.80, 2.74, 3.62]; \\ \mathbf{y} &= [1.01, 0.69, -0.66, -1.34, -1.75, -0.98, 0.25, 1.57, 1.65, 1.51] \end{aligned}$$

We would like to learn a Bayesian linear regression model using this data, assuming a Gaussian likelihood model for the outputs with fixed noise precision  $\beta = 4$ . However, instead of working with the original scalar-valued

inputs, we will map each input  $x$  using a degree- $k$  polynomial as  $\phi_k(x) = [1, x, x^2, \dots, x^k]^\top$ . Note that, when using the mapping  $\phi_k$ , each original input becomes  $k + 1$  dimensional. Denote the entire set of mapped inputs as  $\phi_k(\mathbf{x})$ , a  $10 \times (k + 1)$  matrix. Consider  $k = 1, 2, 3$  and  $4$ , and learn a Bayesian linear regression model for each case. Assume the following prior on the regression weights:  $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{I})$  with  $\mathbf{w} \in \mathbb{R}^{k+1}$ .

1. For each  $k$ , compute the posterior of  $\mathbf{w}$  and show a plot with 10 random functions drawn from the inferred posterior (show the functions for the input range  $x \in [-4, 4]$ ). Also show the original training examples on the same plot to illustrate how well the functions fit the training data.
2. For each  $k$ , compute and plot the **mean** of the posterior predictive  $p(y_*|\phi_k(x_*), \phi_k(\mathbf{x}), \mathbf{y}, \beta)$  on the interval  $x_* \in [-4, 4]$ . On the same plot, also show the predictive posterior mean plus-and-minus two times the predictive posterior standard deviation.
3. Compute the log marginal likelihood  $\log p(\mathbf{y} | \phi_k(\mathbf{x}), \beta)$  of the training data for each of the 4 mappings  $k = 1, 2, 3$ . Which of these 4 “models” seems to explain the data the best?
4. Using the MAP estimate  $\mathbf{w}_{MAP}$ , Compute the log likelihood  $\log p(\mathbf{y}|\mathbf{w}_{MAP}, \phi_k(\mathbf{x}), \beta)$  for each  $k$ . Which of these 4 models seems to have the highest log likelihood? Is your answer the same as that based on the log marginal likelihood (part 3)? Which of these two criteria (highest log likelihood or highest log marginal likelihood) do you think is more reasonable to select the best model and why?
5. For your best model, suppose you could include an additional training input  $x'$  (along with its output  $y'$ ) to “improve” your learned model using this additional example. Where in the region  $x \in [-4, 4]$  would you like the chosen  $x'$  to be? Explain your answer briefly,

You may use MATLAB or Python for this part but you should implement the code yourself (and not use an existing implementation of Bayesian linear regression). Submit the plots as well as the code in a single zip file (named `yourrollnumber.zip`).

## Some Additional Practice Problems (not for credit)

### Problem 1

**(Conjugate Prior for Uniform Distribution)** Consider  $N$  observations  $\{x_1, x_2, \dots, x_N\}$  drawn i.i.d. from Uniform( $0, \theta$ ) distribution, i.e.,  $p(x_n|\theta) = \frac{1}{\theta} \mathbb{I}[0 < x_n < \theta]$ , where  $\mathbb{I}[\cdot]$  denotes the indicator function (equals 1 if the condition is true; and 0 otherwise). Assume a Pareto distribution as the prior on  $\theta$ . The Pareto distribution has the form  $p(\theta|\alpha, c) = \frac{\alpha c^\alpha}{\theta^{\alpha+1}} \mathbb{I}[\theta > c]$ . Derive the posterior distribution of  $\theta$  and use the derivation to show that the Pareto distribution is conjugate to the Uniform distribution.

### Problem 2

**(When You Integrate Out..)** Suppose  $x$  is a scalar random variable drawn from a univariate Gaussian  $p(x|\eta) = \mathcal{N}(x|0, \eta)$ . The variance  $\eta$  itself is drawn from an exponential distribution:  $p(\eta|\gamma) = \text{Exp}(\eta|\gamma^2/2)$ , where  $\gamma > 0$ . Note that the exponential distribution is defined as  $\text{Exp}(x|\lambda) = \lambda \exp(-\lambda x)$ . Derive the expression of the marginal distribution of  $x$ , i.e.,  $p(x|\gamma) = \int p(x|\eta)p(\eta|\gamma)d\eta$  after integrating out  $\eta$ . Plot both  $p(x|\eta)$  and  $p(x|\gamma)$ . What difference do you see between the shapes of these two distributions? **Note:** You don’t need to submit the code used to generate the plots. Just the plots (appropriately labeled) are fine.

**Hint:** You will notice that  $\int p(x|\eta)p(\eta|\gamma)d\eta$  is a hard to compute integral. However, the solution does have a closed form expression. One way to get the result is to compute the **moment generating function (MGF)**<sup>1</sup>

<sup>1</sup>MGF of a p.d.f.  $p(x)$  is defined as  $M_X(t) = \int_{-\infty}^{\infty} e^{tx} p(x) dx$

of  $\int p(x|\eta)p(\eta|\gamma)d\eta$  (note that this is a p.d.f.) and compare the obtained MGF expression with the MGFs of various p.d.f.s given in the table on the following Wikipedia page: [https://en.wikipedia.org/wiki/Moment-generating\\_function](https://en.wikipedia.org/wiki/Moment-generating_function), and identify which p.d.f.'s MGF it matches with. That will give you the form of distribution  $p(x|\gamma)$ . Specifically, name this distribution and identify its parameters.

### Problem 3

**(Integrating Out Rate Parameter of Poisson)** Assume a random variable  $x$  with distribution  $p(x|\lambda) = \text{Poisson}(x|\lambda)$  where  $\lambda > 0$  denotes the rate parameter of the Poisson. Assume the Poisson rate to have a distribution  $p(\lambda|a, b) = \text{Gamma}(a, b)$  with non-negative shape parameter  $a = r$  and rate parameter  $b = (1 - p)/p$  with  $p \in (0, 1)$ . Integrate out  $\lambda$  and give the complete expression for the marginal distribution of  $x$ , i.e.,  $p(x|a, b) = \int p(x|\lambda)p(\lambda|a, b)d\lambda$ . Do you know the name of this distribution? If not, maybe try looking up (this distribution is also used to model counts), and read up about its properties as compared to Poisson.

### Problem 4

**(It Gets Better..)** Recall that, for a Bayesian linear regression model with likelihood  $p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \mathbf{x}, \beta^{-1})$  and prior  $p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1}\mathbf{I})$ , the *predictive posterior* is  $p(y_*|\mathbf{x}_*) = \mathcal{N}(\mu_N^\top \mathbf{x}_*, \beta^{-1} + \mathbf{x}_*^\top \Sigma_N \mathbf{x}_*) = \mathcal{N}(\mu_N^\top \mathbf{x}_*, \sigma_N^2(\mathbf{x}_*))$ , where we have defined  $\sigma_N^2(\mathbf{x}_*) = \beta^{-1} + \mathbf{x}_*^\top \Sigma_N \mathbf{x}_*$  and  $\mu_N$  and  $\Sigma_N$  are the mean and covariance matrix of the Gaussian posterior on  $\mathbf{w}$ , s.t.,  $\mu_N = \Sigma(\beta \sum_{n=1}^N y_n \mathbf{x}_n)$  and  $\Sigma_N = (\beta \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^\top + \lambda \mathbf{I})^{-1}$ . Here, we have used the subscript  $N$  to denote that the model is learned using  $N$  training examples. Show that, as the training set size  $N$  increases, the variance of the predictive posterior goes down, i.e., the model becomes more and more certain about the predictions on the test example  $\mathbf{x}_*$ , i.e.,

$$\sigma_{N+1}^2(\mathbf{x}_*) < \sigma_N^2(\mathbf{x}_*)$$

You may make use the following matrix identity:

$$(\mathbf{M} + \mathbf{v}\mathbf{v}^\top)^{-1} = \mathbf{M}^{-1} - \frac{(\mathbf{M}^{-1}\mathbf{v})(\mathbf{v}^\top \mathbf{M}^{-1})}{1 + \mathbf{v}^\top \mathbf{M}^{-1}\mathbf{v}}$$

Where  $\mathbf{M}$  denotes a square matrix and  $\mathbf{v}$  denotes a column vector.

### Problem 5

**(A Scaled Prior)** Consider a prior  $p(\theta|\alpha)$  conjugate to a likelihood  $p(\mathbf{x}|\theta)$ . Suppose we scale  $p(\theta)$  by multiplying it by another non-negative function  $f(\theta)$  and normalize it (so that it still is a probability distribution) to define another distribution (which will still be defined by parameters  $\alpha$ )

$$q(\theta|\alpha) = \frac{p(\theta|\alpha)f(\theta)}{\int p(\theta|\alpha)f(\theta)d\theta}$$

Assume  $0 < \int p(\theta|\alpha)f(\theta) < \infty$ . Is the new distribution  $q(\theta|\alpha)$  still conjugate to the likelihood  $p(\mathbf{x}|\theta)$ ? Support your answer by showing proper steps.

### Problem 6

**(Normal-Gamma as Exponential Family Distribution)** The normal-gamma distribution defines a joint distribution over two random variables  $\mu \in (-\infty, \infty)$  and  $\tau \in (0, \infty)$  and is defined as

$$p(\mu, \tau|\mu_0, \lambda_0, \alpha_0, \beta_0) = \frac{\beta_0^{\alpha_0} \sqrt{\lambda_0}}{\Gamma(\alpha_0) \sqrt{2\pi}} \tau^{\alpha_0 - \frac{1}{2}} \exp(-\beta_0 \tau) \exp\left(-\frac{\lambda_0 \tau (\mu - \mu_0)^2}{2}\right)$$

The normal-gamma distribution is defined by four parameters  $\mu_0 \in \mathbb{R}$ ,  $\lambda_0 \in \mathbb{R}_+$ ,  $\alpha_0 \in \mathbb{R}_+$ , and  $\beta_0 \in \mathbb{R}_+$ .

Express the normal-gamma distribution in form of an exponential family distribution, and identify all the relevant quantities such as sufficient statistics, natural parameters, log partition function.